

Chapter-1

1) Driver types (Character / Block / Network) — quick summary & when to use

- **Character drivers**

- Handle stream-oriented, unstructured byte I/O (e.g., /dev/mychar).
- Implement file_operations (open, read, write, llseek, unlocked_ioctl, poll/fasync, release).
- No buffering by the block layer by default — driver implements its own buffering or passes data directly.
- Use when access is simple, control-like devices (GPIO, UART, simple sensors).

- **Block drivers**

- Present a block device interface (e.g., disks, eMMC). Implement gendisk, request_queue, blk_mq or older request API.
- Provide random access, sector-based reads/writes, caching, and support filesystem mounting.
- Use when device stores persistent data or must behave like a disk.

- **Network drivers**

- Implement struct net_device operations (open, stop, start_xmit, ndo_get_stats, ethtool ops).
 - Interact with NAPI, interrupts/softirqs, and DMA.
 - Use when device communicates packets (Ethernet, WiFi, CAN).
-

Overview: What Are “Driver Types”?

1. Character drivers

2. Block drivers
3. Network drivers

1 Character Drivers

◆ Concept

- Character drivers deal with data streams — bytes flowing sequentially.
- They allow read/write like a file (no concept of sectors or packets).
- The kernel doesn't buffer or schedule I/O in blocks — it's the driver's job to handle data flow directly.

◆ Example Devices

Example Device	Description	Device Node
UART / Serial port	Byte stream (characters sent one by one)	/dev/ttyS0
GPIO	Simple ON/OFF or value control	/dev/gpiochip0
I2C/SPI sensors	You read sensor values via commands	/dev/i2c-1
ADC/DAC	Continuous analog-to-digital or digital-to-analog conversion	/dev/adc0
Custom user device (e.g., LED control)	You toggle LEDs using IOCTLs	/dev/myled

◆ Key Kernel Structures & Functions

- **Major/Minor numbers** → identify the device
- **file_operations** struct:
- struct file_operations fops = {

- .open = my_open,
- .read = my_read,
- .write = my_write,
- .unlocked_ioctl = my_ioctl,
- .release = my_close,
- };
- Register with:
register_chrdev_region() or alloc_chrdev_region()

◆ When to Choose Character Driver

✓ When your device:

- Transfers data as bytes or streams
- Doesn't need a filesystem
- Doesn't need random access
- Is simple control/stream device (UART, I2C, GPIO, sensors)

2 Block Drivers

◆ Concept

- Block drivers deal with fixed-size blocks or sectors (typically 512B, 4KB).
- The kernel's block layer handles buffering, caching, and request scheduling.

- Filesystems (ext4, FAT, etc.) sit on top of block devices.

◆ Example Devices

Example Device Description		Device Node
Hard Disk	Persistent storage	/dev/sda
SD/eMMC	Flash-based block storage	/dev/mmcblk0
USB Pendrive	External block storage	/dev/sdb
RAM Disk	Memory treated as disk	/dev/ram0

◆ Key Kernel Structures & Functions

- **gendisk** structure — represents a disk
- **request_queue** — handles read/write requests
- **blk-mq (Block Multi-Queue)** — modern high-performance API
- Register with:
 `add_disk()` after initializing gendisk structure

Example (simplified):

```
struct gendisk *gd = alloc_disk(1);
gd->major = major;
gd->fops = &block_fops;
gd->queue = blk_init_queue(my_request_fn, &lock);
```

```
add_disk(gd);
```

◆ When to Choose Block Driver

✓ When your device:

- Stores persistent data
- Needs random (seek-based) access
- Should support filesystem mounting
- Acts like a disk or flash memory

Example Decision:

- NAND Flash (raw): Character driver
 - eMMC / SD card: Block driver (supports filesystem mount)
-

3 Network Drivers

◆ Concept

- Network drivers send and receive packets, not bytes or blocks.
- The Linux network stack (TCP/IP layers) sits above them.
- They manage struct net_device, DMA rings, interrupts, and packet queues.

◆ Example Devices

Example Device	Description	Interface Name
Ethernet Controller	Wired network	eth0
Wi-Fi Module	Wireless network	wlan0
CAN Controller	Automotive network	can0
Virtual Network	Software device	tun0, tap0

◆ Key Kernel Structures & Functions

- `struct net_device` → represents a network interface
- `net_device_ops`:
- `static const struct net_device_ops my_netdev_ops = {`
- `.ndo_open = my_open,`
- `.ndo_stop = my_close,`
- `.ndo_start_xmit = my_xmit,`
- `.ndo_get_stats = my_stats,`
- `};`
- Register with: `register_netdev(ndev);`

◆ When to Choose Network Driver

✓ When your device:

- Sends/receives network packets

- Connects to a network (Ethernet, WiFi, CAN)
- Works with the kernel's TCP/IP stack

How to Decide Which Driver Type to Use

Decision Criteria	Choose Character Driver	Choose Block Driver	Choose Network Driver
Data type	Stream / byte-based	Block / sector-based	Packet-based
Access mode	Sequential	Random (seekable)	Network frame
Example	UART, I2C, GPIO	HDD, eMMC, SD card	Ethernet, WiFi, CAN
Kernel Interface	/dev/mydevice	/dev/sda, /dev/mmcblk0	eth0, wlan0
Filesystem mount	✗ No	✓ Yes	✗ No
Typical operations	read/write/ioctl	read/write sectors	send/receive packets

Quick Real-Time Examples by Use Case

Use Case	Example Device	Driver Type	Reason
LED driver	GPIO output control	Character	Simple control (on/off)
Temperature sensor (I2C)	LM75	Character	Byte read/write commands

Use Case	Example Device	Driver Type	Reason
eMMC storage	SD/eMMC	Block	Sector-based persistent storage
Ethernet controller	LAN8720	Network	Transmit/receive packets
USB mass storage	Pendrive	Block	Acts like a disk
UART-based GPS	GPS receiver	Character	Continuous data stream
Wi-Fi module	RTL8188	Network	Packet-based communication

In Real Embedded Project — Example Scenario

Let's say you're building an **Embedded Linux-based Smart TV**:

Hardware Component	Driver Type	Why
eMMC storage	Block	For filesystem & media files
Wi-Fi module	Network	For internet connection
Bluetooth chip	Character	For serial control
IR Remote sensor	Character	For decoding signals
HDMI output	Character (or DRM driver)	Control interface

Hardware Component	Driver Type	Why
LED status	Character	Simple control on/off
Audio Codec (I2S)	Character	Stream of samples

Interview tip: name examples, and explain why you chose each type (e.g., "character for a simple sensor that needs ioctl commands").

2) Major/minor numbers & register_chrdev_region / alloc_chrdev_region / cdev

- **Device number (dev_t)** = major (driver identity) + minor (per-device instance).
 - Major: which driver handles the device numbers.
 - Minor: device instance (e.g., /dev/mydev0, /dev/mydev1).
- **Static vs dynamic allocation**
 - register_chrdev_region(dev_t first, unsigned count, const char *name) — static major (you know the major number).
 - alloc_chrdev_region(dev_t *dev, unsigned baseminor, unsigned count, const char *name) — dynamic major (recommended).
- **char device registration pattern**

```
dev_t dev;
```

```
struct cdev my_cdev;
```

```
int rc;
```

```
/* allocate dynamic major */
rc = alloc_chrdev_region(&dev, 0, 1, "mychardev");
if (rc) return rc;
```

```
/* init cdev and add */
cdev_init(&my_cdev, &my_fops);
my_cdev.owner = THIS_MODULE;
rc = cdev_add(&my_cdev, dev, 1);
if (rc) {
    unregister_chrdev_region(dev, 1);
    return rc;
}
```

- **Unload cleanup**

```
cdev_del(&my_cdev);
unregister_chrdev_region(dev, 1);
```

Major & Minor Numbers — Theory

- ◆ **What Are They?**

In Linux, every device (whether it's `/dev/ttyS0`, `/dev/sda`, or `/dev/mydevice`) is represented as a special file inside `/dev/`.

The kernel uses major and minor numbers to map that file to the correct driver.

Number Purpose

Major number Identifies the **driver** (which kernel module handles this device)

Minor number Identifies the **specific device instance** the driver controls

◆ Real Example

Let's say you create two LED devices:

`/dev/led0`

`/dev/led1`

They might share the same major number (since they use the same driver), but have different minor numbers:

Device File	Major	Minor
<code>/dev/led0</code>	240	0
<code>/dev/led1</code>	240	1

This means:

- **Major 240** → handled by **your LED driver**.
- **Minor 0/1** → tells your driver which LED instance to access.

3 Ways to Register Device Numbers

There are **two main APIs** for registering device numbers (and one for dynamic management):

1 register_chrdev_region() — Static Registration

- You manually choose a major number.
- Used when you want fixed numbers (e.g., for known hardware or system drivers).

Prototype:

```
int register_chrdev_region(dev_t dev, unsigned count, const char *name);
```

Example:

```
dev_t devno = MKDEV(240, 0);
register_chrdev_region(devno, 2, "my_led_driver");
```

Major = 240, Minor = 0 to 1 registered for 2 devices.

Used when you know the major number and want consistent / predictable device nodes.

2 alloc_chrdev_region() — Dynamic Registration

- Kernel dynamically allocates a **free major number** for you.
- Common for **loadable kernel modules** (LKMs).

Prototype:

```
int alloc_chrdev_region(dev_t *dev, unsigned baseminor, unsigned count, const char *name);
```

Example:

```
dev_t devno;  
alloc_chrdev_region(&devno, 0, 2, "my_led_driver");  
int major = MAJOR(devno);  
int minor = MINOR(devno);
```

The kernel assigns an available **major number**.

You can print it using `printk()` to see which one was allocated.

Used in modern drivers for flexibility and to avoid conflicts.

3 unregister_chrdev_region() — Cleanup

- Always free your allocated region in the exit function:

```
unregister_chrdev_region(devno, 2);
```

4 The cdev Structure — Connecting to File Operations

Once your major/minor numbers are registered, you must link them to your driver's functionality.

This is done using a `struct cdev` (Character Device Structure).

struct cdev

Represents your character device in the kernel — holds:

- File operation methods (read, write, ioctl, etc.)
 - The device numbers (major/minor)
 - Ownership information
-

Steps to Initialize and Add cdev

Example:

```
struct cdev my_cdev;  
cdev_init(&my_cdev, &fops); // Link with file operations  
cdev_add(&my_cdev, devno, 2); // Add to kernel
```

Cleanup:

```
cdev_del(&my_cdev);
```

Full Initialization Flow (Summary)

Step	API	Purpose
1	alloc_chrdev_region()	Get major/minor numbers
2	cdev_init()	Initialize your cdev structure
3	cdev_add()	Register your device with kernel

Step API**Purpose**

- | | |
|--|--|
| 4 class_create() + device_create() | Create /dev/mydevice node (optional for auto udev) |
| 5 cdev_del() + unregister_chrdev_region() | Cleanup during exit |
-

Example Code (Typical Template)

```
#include <linux/module.h>
```

```
#include <linux/fs.h>
```

```
#include <linux/cdev.h>
```

```
#define DEVICE_COUNT 1
```

```
static dev_t devno;
```

```
static struct cdev my_cdev;
```

```
static struct class *cls;
```

```
static struct file_operations fops = {
```

```
    .owner = THIS_MODULE,
```

```
    .open = my_open,
```

```
.read = my_read,  
.write = my_write,  
.release = my_close,  
};  
  
static int __init my_init(void)  
{  
    // Step 1: Allocate device numbers  
    alloc_chrdev_region(&devno, 0, DEVICE_COUNT, "my_char_driver");  
    printk("Major: %d Minor: %d\n", MAJOR(devno), MINOR(devno));  
  
    // Step 2: Init and add cdev  
    cdev_init(&my_cdev, &fops);  
    cdev_add(&my_cdev, devno, DEVICE_COUNT);  
  
    // Step 3: Create device node automatically (optional)  
    cls = class_create(THIS_MODULE, "my_class");  
    device_create(cls, NULL, devno, NULL, "my_char_device");  
}
```

```
    printk("Character driver registered successfully\n");
    return 0;
}

static void __exit my_exit(void)
{
    device_destroy(cls, devno);
    class_destroy(cls);
    cdev_del(&my_cdev);
    unregister_chrdev_region(devno, DEVICE_COUNT);
    printk("Character driver unregistered\n");
}

module_init(my_init);
module_exit(my_exit);
MODULE_LICENSE("GPL");
```

Interview-Style Summary Answer

“In Linux, every device is identified by a **major** and **minor** number.

The major number tells the kernel which driver to call, and the minor number identifies the specific device handled by that driver.

For character drivers, I register these numbers using `alloc_chrdev_region()` or `register_chrdev_region()`.

Then, I link my driver’s file operations using a struct `cdev`, initialized with `cdev_init()` and added via `cdev_add()`.

When unloading, I clean up using `cdev_del()` and `unregister_chrdev_region()`.

I usually prefer `alloc_chrdev_region()` for dynamic major number allocation to avoid conflicts.”

3) file_operations: open, read, write, ioctl (signatures & behavior)

- Important file_operations members:

```
int (*open)(struct inode *inode, struct file *file);
```

```
ssize_t (*read)(struct file *file, char __user *buf, size_t count, loff_t *ppos);
```

```
ssize_t (*write)(struct file *file, const char __user *buf, size_t count, loff_t *ppos);
```

```
long (*unlocked_ioctl)(struct file *file, unsigned int cmd, unsigned long arg);
```

```
int (*release)(struct inode *inode, struct file *file);
```

```
unsigned int (*poll)(struct file *file, struct poll_table_struct *wait);
```

- Use `container_of(inode->i_cdev, struct mydev, cdev)` in `open` to attach driver private data to `file->private_data` for subsequent operations.
- **IOCTL** for control/complex config: implement with `copy_from_user` if passing a struct pointer. Prefer `ioctl` for non-stream control operations.

Example: open/read pattern

```
static int my_open(struct inode *inode, struct file *file) {
    struct mydev *dev = container_of(inode->i_cdev, struct mydev, cdev);
    file->private_data = dev;
    return 0;
}
```

4) copy_to_user / copy_from_user & get_user / put_user

- Kernel cannot directly access userspace pointers — must use safe accessors:
 - copy_to_user(void __user *to, const void *from, unsigned long n) returns number of bytes not copied (0 = success).
 - copy_from_user(void *to, const void __user *from, unsigned long n) same semantic.
 - For scalar values, prefer get_user(x, ptr) / put_user(x, ptr) (returns 0 on success).
- **Always check return** and return -EFAULT on failure.
- Example read:

```
ssize_t my_read(struct file *file, char __user *buf, size_t count, loff_t *ppos) {
    struct mydev *d = file->private_data;
    size_t avail = d->buf_len - *ppos;
    size_t to_copy = min(avail, count);
    if (to_copy == 0) return 0;
    if (copy_to_user(buf, d->kbuf + *ppos, to_copy)) return -EFAULT;
```

```

*ppos += to_copy;
return to_copy;
}

```

5) Wait queues & poll/select (non-blocking and blocking I/O)

- **Wait queues** let processes sleep until condition becomes true.
 - DECLARE_WAIT_QUEUE_HEAD(wq);
 - wait_event_interruptible(wq, condition); — sleeps until condition true or signal; returns -ERESTARTSYS on signal.
 - wake_up_interruptible(&wq) wakes waiters.
- Typical pattern for read():

```

if (buffer_empty) {
if (file->f_flags & O_NONBLOCK)
return -EAGAIN;
rc = wait_event_interruptible(dev->wq, data_available());
if (rc) return rc; // signal interrupted
}

/* now read data */

```

- **poll/select** support via poll() in file_operations:

```

unsigned int my_poll(struct file *file, poll_table *wait) {

```

```

struct mydev *dev = file->private_data;
poll_wait(file, &dev->wq, wait);
unsigned int mask = 0;
if (data_available()) mask |= POLLIN | POLLRDNORM;
if (space_available()) mask |= POLLOUT | POLLWRNORM;
return mask;
}

```

- `poll_wait` registers the wait queue so `select/poll` knows where to sleep; when your code does `wake_up_interruptible`, pollers are notified.

Interview tip: highlight non-blocking behaviour and signal handling; mention `fasync/kill_fasync` if supporting `SIGIO`.

6) Synchronization — spinlocks, semaphores, `atomic_t`, mutexes

- **Spinlocks (`spinlock_t`)**
 - Use in interrupt context or short critical sections; do not sleep.
 - `spin_lock_irqsave(&lock, flags); /* critical */ spin_unlock_irqrestore(&lock, flags);`
 - Use for protecting small data structures where blocking is unacceptable.
- **Mutex (`struct mutex`)**
 - Sleeping lock for process context only; use when you may sleep (e.g., IO wait).
 - `mutex_init(&m); mutex_lock(&m); mutex_unlock(&m);`
- **Semaphores (`struct semaphore`)**

- Older API; similar to mutex but counting; `down_interruptible`, `down_trylock`, `up`.
- **Atomic operations (`atomic_t`)**
 - For simple counters/flags where only atomic inc/dec/test is needed.
 - `atomic_t cnt = ATOMIC_INIT(0); atomic_inc(&cnt); if (atomic_dec_and_test(&cnt)) ...`
- **Which to use?**
 - Short non-sleeping ops -> spinlock.
 - Might sleep or blocking IO -> mutex.
 - Simple counters -> `atomic_t`.

◆ What is Synchronization?

In the Linux kernel, multiple threads or CPUs can access the same shared resource (e.g., a global variable, hardware register, buffer). If two or more threads modify it at the same time, unpredictable behavior occurs.

This is called a Race Condition.

⚠ **Race Condition** — Example

Code Example:

```
int counter = 0;

void write_func(void)
{
    counter++;
}
```

```
}
```

If two threads call `write_func()` simultaneously on multi-core CPU:

1. Both read counter = 0
2. Both increment it to 1
3. Both write back 1

✗ Expected = 2

✓ Actual = 1

Data corruption or inconsistent state — this is a race condition.

Why Synchronization Is Needed?

To prevent race conditions and ensure mutual exclusive access to shared data.

In Linux kernel, we use synchronization primitives like:

- Spinlocks
 - Semaphores
 - Mutexes
 - Atomic operations
-

1 Spinlocks

Concept

- Spinlock is a busy-wait lock.
- The thread spins (loops) until the lock becomes free.
- It never sleeps, so it is used in interrupt context or short critical sections.

Key Points

Property	Description
Sleeping allowed?	✗ No (used in atomic/interrupt context)
Suitable for	Short, fast critical sections
Blocking behavior	Spins until lock available
Reentrant?	✗ No
IRQ safe?	Use <code>spin_lock_irqsave()</code> if in interrupt context

Example:

```
spinlock_t my_lock;
spin_lock_init(&my_lock);
void my_function(void)
{
    spin_lock(&my_lock);
```

```
// critical section  
counter++;  
spin_unlock(&my_lock);  
}
```

In interrupt context:

```
unsigned long flags;  
spin_lock_irqsave(&my_lock, flags);  
// critical section  
spin_unlock_irqrestore(&my_lock, flags);
```

When to Use:

✓ Use spinlock when:

- Code runs in interrupt context
- You cannot sleep
- Critical section is very short
- Used for protecting registers, flags, or shared buffers

Example:

Protecting a shared ring buffer between interrupt handler and bottom half (tasklet/workqueue).

♦ **2** Semaphores

Concept

- Semaphore allows limited concurrent access (counting semaphore) or mutual exclusion (binary semaphore).
- Unlike spinlocks, it can sleep, so it's suitable for process context only.

Key Points

Property	Description
Sleeping allowed?	✔ Yes
Type	Counting or Binary
Suitable for	Long operations
Context	Process context only (not interrupt)

Example:

```
struct semaphore my_sem;
sema_init(&my_sem, 1); // Binary semaphore
void my_function(void)
{
    down(&my_sem); // acquire (may sleep)
```

```
// critical section  
counter++;  
up(&my_sem); // release  
}
```

When to Use:

Use semaphore when:

- Critical section can block or sleep
- Protect shared resource accessed by multiple processes
- Operation may take long time

Example:

Synchronizing between read() and write() in a character driver where user-space may sleep.

3 Mutex

Concept

- Mutex = “Mutual Exclusion Lock”
- Similar to a binary semaphore, but designed only for mutual exclusion (not counting).
- Simpler & faster for 1-thread-at-a-time protection.
- Cannot be used in interrupt context.

Key Points

Property	Description
Sleeping allowed?	✓ Yes
Recursive locking allowed?	✗ No
Suitable for	Long critical section, process context
API	mutex_lock() / mutex_unlock()

Example:

```
struct mutex my_mutex;
mutex_init(&my_mutex);
void my_function(void)
{
    mutex_lock(&my_mutex);
    // critical section
    counter++;
    mutex_unlock(&my_mutex);
}
```

When to Use:

Use mutex when:

- Only one thread should access a resource
- Operation can sleep
- Code runs in process context

Example:

File operations (open, read, write) in a driver where only one user should access hardware at a time.

Atomic Operations

Concept

- Used for simple integer operations (increment, decrement, set) that must happen atomically.
 - Very fast, no locking required.
 - Usually implemented using CPU atomic instructions.
-

Key Points

Property	Description
----------	-------------

Sleeping allowed?	 Yes
-------------------	---

Property	Description
Type	Integer / bit operations
Locking needed?	✗ No
Performance	Fastest

Example:

```
atomic_t counter = ATOMIC_INIT(0);
```

```
void func(void)
```

```
{  
    atomic_inc(&counter);  
    atomic_dec(&counter);  
    int val = atomic_read(&counter);  
}
```

Or for bit operations:

```
set_bit(bit, &flags);
```

```
clear_bit(bit, &flags);
```

```
test_and_set_bit(bit, &flags);
```

When to Use:

Use atomic_t when:

- You only need atomic integer or bit modification
- Critical section is just one variable
- No complex logic or sleeping needed

Example:

Counting the number of active open file handles or reference counts.

Comparison Summary Table

Feature	Spinlock	Semaphore	Mutex	Atomic_t
Sleep allowed	✗ No	✓ Yes	✓ Yes	✓ Yes
Context	Interrupt / process	Process only	Process only	Both
Used for	Short atomic section	Long blocking ops	Mutual exclusion	Simple counters
Kernel wait behavior	Busy wait	Sleep	Sleep	Non-blocking
API Example	spin_lock()	down() / up()	mutex_lock()	atomic_inc()

Typical Real-Time Use Cases

Scenario	Recommended Primitive Reason	
Protect shared buffer between interrupt & process	<code>spin_lock_irqsave()</code>	Cannot sleep in interrupt
Protect global variable between two user-space operations	<code>mutex</code>	Can sleep
Protect I2C device access (long hardware access)	<code>semaphore</code>	Blocking allowed
Count number of open files	<code>atomic_t</code>	Just integer operation

Interview Answer Example (2-Minute Version)

“Synchronization is needed to prevent race conditions when multiple contexts access shared data.

If a driver’s interrupt handler and process context share a variable, a race condition may occur.

To avoid this, I use different kernel primitives based on context:

- Spinlock — for short critical sections in interrupt or atomic context (no sleeping).
- Semaphore — when multiple processes access a shared resource, and the code can sleep.
- Mutex — for simple one-at-a-time access in process context.
- Atomic_t — for fast atomic operations like counters or flags.

For example, I would use a spinlock to protect a shared buffer between ISR and bottom half, and a mutex in read()/write() to ensure exclusive user-space access.”

7) Workqueues, tasklets, softirqs — bottom-half mechanisms & differences

- **Context hierarchy**

- Top half = ISR (interrupt context) — minimal work, must not sleep.
- Bottom halves = deferred work that runs later to finish processing.

- **Softirq**

- Low-level kernel mechanism; per-CPU; cannot sleep; used by networking, timers.
- Scheduled with `raise_softirq` or `open_softirq`. Hard to use directly; use tasklets or workqueues instead.

- **Tasklet**

- Built on top of softirq; per-CPU, runs in softirq context, cannot sleep.
- Use `DECLARE_TASKLET` or `tasklet_init(&t, func, data); tasklet_schedule(&t);`
- Serialized execution for same tasklet on single CPU; not safe to sleep.

- **Workqueue (recommended)**

- Work runs in process context in kernel threads — **may sleep**.
- Use `DECLARE_WORK(&work, work_fn); schedule_work(&work);` or create custom workqueue with `alloc_workqueue`.
- Preferred for complex or blocking deferred work.

- **When to use which?**

- Minimal, per-CPU, fast handlers -> softirq/tasklet.
- Needs to sleep, do blocking I/O, or long work -> workqueue.

- **Example: ISR schedules workqueue**

```
static irqreturn_t my_isr(int irq, void *dev_id) {
    struct mydev *dev = dev_id;

    /* Do minimal check and queue work */
    schedule_work(&dev->work);

    return IRQ_HANDLED;
}
```

```
static void my_work_fn(struct work_struct *work) {
    struct mydev *dev = container_of(work, struct mydev, work);

    /* heavy processing allowed, can sleep */
}
```

Interview tip: they often test the difference: “*Which one can sleep?*” — answer: workqueue can, tasklet/softirq cannot.

8) Platform driver model & of_match_table

- **Platform device / driver concept**
 - platform_device represents a device that’s not discoverable by bus (e.g., SoC peripherals). Registered either by board code or by device tree (/proc/device-tree).
 - platform_driver contains probe and remove functions and driver struct (with name and of_match_table).
- **Matching flow**

- Kernel matches device to driver by:
 - Device tree compatible property vs driver of `_match_table`, OR
 - `name / id_table` for non-DT platform devices

- **Typical driver skeleton**

```
static int my_probe(struct platform_device *pdev) {
    struct resource *res;

    res = platform_get_resource(pdev, IORESOURCE_MEM, 0);
    devm_ioremap_resource(&pdev->dev, res);

    /* get IRQ */

    irq = platform_get_irq(pdev, 0);
    request_irq(irq, my_isr, 0, dev_name(&pdev->dev), dev);

    /* setup, devm allocations, register char device, etc. */

    return 0;
}

static int my_remove(struct platform_device *pdev) {
    /* free resources are automatically cleaned if devm used, otherwise free explicitly */

    return 0;
}
```

```
static const struct of_device_id my_of_match[] = {
    { .compatible = "vendor,mydev", },
    {},
};
MODULE_DEVICE_TABLE(of, my_of_match);
```

```
static struct platform_driver my_driver = {
    .probe = my_probe,
    .remove = my_remove,
    .driver = {
        .name = "my_driver",
        .of_match_table = of_match_ptr(my_of_match),
    },
};
module_platform_driver(my_driver);
```

- **Best practices**

- Use devm_ APIs (devm_kzalloc, devm_ioremap_resource, devm_request_threaded_irq) to simplify cleanup.
- Read DT properties with of_property_read_u32, of_find_property, or devm_kstrdup.

Interview tip: explain DT compatible usage and how you would debug DT matching (check dmesg for "of_match" logs, cat /proc/device-tree/...).

9) Sysfs vs procfs (differences and usage)

- **sysfs**

- Exposes kernel objects (kobject, device, driver) as attribute files (structured hierarchy under /sys).
- Per-device attributes created by `device_create_file()` or `sysfs_create_group()`.
- Use for exposing device attributes and controls (state, configuration).
- Attributes use show and store:

```
static ssize_t foo_show(struct device *dev, struct device_attribute *attr, char *buf) {
    return sprintf(buf, "%d\n", my_val);
}

static ssize_t foo_store(struct device *dev, struct device_attribute *attr, const char *buf, size_t count) {
    sscanf(buf, "%d", &my_val);
    return count;
}

DEVICE_ATTR(myattr, 0644, foo_show, foo_store);
device_create_file(dev, &dev_attr_myattr);
```

- **procfs**

- General-purpose pseudo-filesystem under /proc, historically for kernel info and process info.
- Use via `proc_create()` and `proc_ops` to expose arbitrary diagnostic info.
- Better: use `debugfs` for debugging interfaces (temporary, can be disabled).
- **Which to use?**
 - Device configuration/state -> **sysfs**.
 - Generic kernel statistics and debug output -> `procfs` or `debugfs` (`debugfs` preferred for dev-only debug).

Interview tip: mention `sysfs` is for structured device attributes tied to device model; `procfs` is more generic.

10) GPIO, I2C, SPI, UART driver flows — practical patterns

GPIO (`gpio-lib`)

- **Modern API:** descriptor-based `gpiod_*` and `devm_gpiod_get`.
- Typical steps:
 1. Get gpio descriptor: `devm_gpiod_get(&pdev->dev, "reset", GPIOD_OUT_LOW);`
 2. Set value: `gpiod_set_value(desc, 1);`
 3. Optionally request IRQ on gpio: `gpiod_to_irq(desc)` then `request_irq`.
 4. Release automatically if `devm` used.
- Older API: `gpio_request`, `gpio_direction_output`, `gpio_set_value`.

I2C

- I2C client driver steps:

1. Define `i2c_driver` with `id_table` and `of_match_table`.
 2. `probe(struct i2c_client *client, const struct i2c_device_id *id)` — use `i2c_smbus_read_byte_data` or `i2c_transfer`.
 3. `i2c_set_clientdata(client, data);`
 4. Remove cleans up.
- Use `devm_kzalloc` for allocation to auto cleanup.

SPI

- SPI master handles low-level clocking; SPI devices are `spi_device`.
- Steps:
 1. Implement `spi_driver` with `probe(struct spi_device *spi)`.
 2. Setup `spi->max_speed_hz`, `mode`, call `spi_setup`.
 3. Transfer via `spi_sync()` or `spi_async()` using `struct spi_message` and `struct spi_transfer`.
- `spi_sync()` is blocking, good for most drivers.

UART (serial core)

- Typically register a `uart_driver` and `uart_port` or implement `serial_core` bindings.
- Implement `struct uart_ops` (`startup`, `stop`, `tx_empty`, `set_termios`, `start_tx`, `stop_tx`).
- For simple character serial, sometimes user-space `tty` and `serial_core` handles protocol.

Interview tip: mention devicetree properties (`gpios`, `i2c` address, `reg` for SPI), and show how to read them inside `probe`.

11) Interrupt handling, bottom halves, `irqreturn_t` return types

- **Request/free IRQ**

```
int request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags,
const char *name, void *dev);

void free_irq(unsigned int irq, void *dev_id);
```

- **ISR signature**

```
irqreturn_t my_isr(int irq, void *dev_id) {
struct mydev *dev = dev_id;

/* Minimal processing */

schedule_work(&dev->work);

return IRQ_HANDLED; /* or IRQ_NONE */
}
```

- **Return values**

- IRQ_NONE — not our interrupt.
- IRQ_HANDLED — interrupt handled.
- IRQ_WAKE_THREAD — used when using request_threaded_irq() — top half returns IRQ_WAKE_THREAD to wake thread handler.

- **Threaded interrupts**

- request_threaded_irq(irq, top, thread_fn, flags, name, dev);
- top can be NULL -> solely threaded handler runs in process context (can sleep).

- **Other helpers**

- `disable_irq`, `enable_irq`, `disable_irq_nosync`, `synchronize_irq`.
- `synchronize_irq` waits for currently running IRQ handler to finish.

Interview tip: if kernel asks about race conditions in ISR, say you should minimize ISR work and use bottom halves or threaded IRQ to handle heavy processing.

Now — Answering the expected interview questions (concise, high-value)

Q1 — *Explain difference between softirq, tasklet, and workqueue*

- **Softirq**

- Low-level kernel mechanism for deferred work.
- Per-CPU; cannot sleep; may run with interrupts enabled.
- Used by core subsystems (networking, timers).

- **Tasklet**

- Built on softirq abstraction; per-CPU; cannot sleep.
- Good for quick deferred jobs; two types of serialization: same tasklet is serialized on a CPU.
- **Limitation:** cannot block/sleep.

- **Workqueue**

- Runs in kernel thread (process context) — **can sleep** and perform blocking operations.
- Supports parallelism and priorities; easier for complex or long-running work.

- **Practical guidance**

- Use **tasklet/softirq** for short, fast work that must run quickly and not block.
- Use **workqueue** for anything that might sleep, access user-space, or perform longer tasks.

Interview answer style:

“Softirq is a low-level per-CPU deferred mechanism; tasklets sit on softirqs and are non-sleepable; workqueues run in kernel threads and can sleep — so for anything blocking use workqueues.”

Q2 — *How do you share data between kernel and user space?*

Common ways (pros/cons + example use cases):

1. **Character device read/write** — simplest: userspace opens `/dev/mydev` and uses `read()/write()`; kernel uses `copy_to_user/copy_from_user`.
 - Good for streaming data, simple control.
2. **IOCTL** — control commands passing small structs or numbers; use `copy_from_user` inside ioctl handler.
 - Good for configuration commands.
3. **sysfs attributes** (`/sys/...`) — read/write small config/status values.
 - Good for exposing device attributes to scripts.
4. **procfs / debugfs** — expose debug info or counters. debugfs often preferred for dev debug data.
5. **mmap** — create shared mapping between user and kernel for high-speed zero-copy I/O (use `remap_pfn_range` or `vm_insert_page`) — more complex and requires careful DMA / page pinning.

6. **Netlink sockets / character devices / relayfs** — used for complex message passing (netlink is common for kernel<->userspace messaging).
7. **udev + sysfs** — for device events & properties.

Important considerations

- Ownership (who frees buffers), locking, synchronization, blocking vs non-blocking, security checks, and copy errors (EFAULT).
- For large/multi-threaded transfers prefer mmap (if safe) or efficient ring buffers with proper synchronization.

Small code demo (read/write) shown earlier with `copy_to_user`.

Interview tip: show you understand tradeoffs: `ioctl` for control vs `mmap` for performance; security implications.

Q3 — *How does probe and remove work in platform drivers?*

Flow:

1. Device creation

- Platform devices are created by board code, ACPI, or device tree (DT nodes).
- DT node has `compatible = "vendor,device"`; `reg = <...>`; `interrupts = <...>`;

2. Driver registration

- When your `platform_driver` is registered by `module_platform_driver()`, the driver core tries to match registered devices with drivers:
 - Match using `of_match_table` (DT compatible) or `id_table/name`.

3. Matching succeeds -> `probe()`

- Kernel calls your probe(struct platform_device *pdev).
- Typical probe tasks:
 - Get resources: platform_get_resource(...) or devm_ioremap_resource.
 - Parse device tree properties: of_property_read_*.
 - Allocate driver data: devm_kzalloc.
 - Request IRQ: devm_request_threaded_irq or request_irq.
 - Initialize hardware, register char device or other kernel interfaces, create sysfs entries.
 - Store pointer via platform_set_drvdata(pdev, data).
- On any failure, probe must undo allocated items and return error code.

4. Remove

- When driver is unloaded or device removed, remove(struct platform_device *pdev) is called.
- Remove must reverse resource acquisition: destroy sysfs, unregister devices, free IRQs (unless devm_ was used, then kernel frees).
- If module is removed, the driver should ensure no in-flight work remains (cancel workqueue, flush pending works).

5. Hotplug & runtime PM

- Drivers should handle dynamic runtime suspend/resume and probe removal under runtime PM.

Interview tip: explain the **error-path cleanup** in probe — allocate devm resources when possible to avoid explicit free in remove and to simplify code.

Q4 — *What happens during insmod and rmmod? (module lifecycle)*

- **insmod mymod.ko**

1. insmod loads the ELF (.ko) into kernel memory via module loader.
2. Resolve symbol dependencies (if unresolved, modprobe would pull needed modules).
3. Kernel maps module sections, applies relocations and sets up module structure.
4. **module_init()** function is called (your .init), e.g., static int __init my_init(void):
 - Often calls platform_driver_register, module_platform_driver, or cdev_add, alloc_chrdev_region.
 - Performs hardware initialization and registers kernel interfaces.
5. If the module registers a device driver, the driver core may immediately call probe() for matching devices (e.g., platform devices).
6. Module reference counting is set so it cannot be removed while in use.

- **rmmod mymod**

1. rmmod requests module removal; kernel checks module refcount. If >0 (in use), removal fails with EBUSY.
2. If safe, kernel calls **module_exit()** function (your cleanup).
 - Must unregister drivers (platform_driver_unregister or cdev_del / unregister_chrdev_region), free IRQs, cancel workqueues, remove sysfs files, free memory.
3. Kernel unmaps and frees module memory.
4. If devices were registered, they are unbound as part of driver unregister.

Key details to say in interview

- Emphasize **order**: register interfaces in init and un-register in exit in reverse order.

- Modules are reference counted; cannot `rmmod` while open file descriptors or other module refs exist.
- Use `devm_` helpers to simplify cleanup and avoid leaks.
- Always ensure no pending work (`flush_work`, `cancel_work_sync`) during exit.

Extra: Short checklist you can verbally walk through in interviews

- For probe: parse DT -> get resources -> request IRQ -> map I/O -> setup character device (`cdev_add`) -> create sysfs -> init workqueue/tasklets -> mark ready.
- For remove: stop device -> disable IRQ -> cancel work -> remove sysfs -> `cdev_del` -> unmap -> free resources.
- For `insmod/rmmod`: mention symbol resolution, `module_init/module_exit`, refcount check, and driver probe/unbind.

Top 10 Mock Interview Questions & Answers (Device Driver Core)

1 What happens when you do `insmod mydriver.ko`?

Concept:

When you insert a module, the kernel loader loads the `.ko` (ELF binary) into kernel memory, resolves dependencies, and runs your `module_init()` function. That function typically registers your driver (e.g., character, platform, or bus driver).

Technical Answer:

When I execute `insmod`, the kernel verifies the module's ELF header, allocates kernel memory for it, and resolves symbol dependencies. Then it calls the module's init function. Inside that, we usually register our driver using `register_chrdev_region`, `cdev_add`, or `platform_driver_register`. If it's a platform driver, after registration the kernel automatically matches and invokes the `probe()` function for compatible devices defined in the device tree.

Bonus tip:

Mention cleanup symmetry — "During `rmmod`, kernel calls `module_exit()` which unregisters everything and frees memory."

2 What's the difference between `register_chrdev_region()` and `alloc_chrdev_region()`?**Concept:**

Both allocate major/minor numbers for a char device.

- `register_chrdev_region()` → you specify major.
- `alloc_chrdev_region()` → kernel dynamically assigns major.

Technical Answer:

`register_chrdev_region()` is used when we already know the major number and want to register it statically.

`alloc_chrdev_region()` asks the kernel to allocate a free major number dynamically.

Dynamic allocation is preferred to avoid conflicts and improve portability.

Bonus tip:

You can say: "After allocation, I use `MAJOR(dev_t)` and `MINOR(dev_t)` macros to get the values and register my cdev."

3 What are `copy_to_user()` and `copy_from_user()`? Why are they needed?**Concept:**

Kernel and user space have separate address spaces. Direct dereferencing of user pointers in kernel leads to page faults.

Technical Answer:

These are safe APIs to transfer data between user and kernel space.

`copy_from_user()` copies data *from user space to kernel space*, and `copy_to_user()` does the reverse.

They also handle page faults safely and return the number of bytes not copied, so I check for non-zero return and handle -EFAULT.

Bonus tip:

Add that you never trust user pointers: "Always validate size and pointer before copying."

4 What is the difference between spinlock, semaphore, and mutex?

Concept:

All are synchronization primitives, used depending on context and whether sleep is allowed.

Technical Answer:

spinlock is a busy-wait lock, used in interrupt context or short non-sleeping critical sections.

mutex is a sleeping lock for process context where sleeping is allowed.

semaphore is similar to a mutex but can allow multiple counts (used less in modern kernels).

We also have `atomic_t` for simple atomic counters.

Bonus tip:

Mention “`spin_lock_irqsave`” if inside ISR: "I disable local interrupts when using spinlocks in interrupt context to avoid deadlocks."

5 What's the difference between tasklet, softirq, and workqueue?

Concept:

All handle deferred work from interrupt context (bottom halves).

Technical Answer:

Softirqs are the lowest level, per-CPU deferred mechanisms that can't sleep and are used by the kernel itself (like networking).

Tasklets are built on top of softirqs — easier to use but still non-sleepable.

Workqueues run in process context via kernel threads, so they can sleep and perform blocking operations like I/O.

Bonus tip:

Say: "If my ISR needs to sleep or handle heavy processing, I always use workqueues instead of tasklets."

6 How do you share data between user space and kernel space?**Concept:**

You can use character devices, sysfs, IOCTLs, mmap, etc.

Technical Answer:

I mostly use character device interfaces where user space reads/writes data through /dev/mydev using `copy_to_user` and `copy_from_user`.

For control operations, I use IOCTLs, and for configuration parameters, sysfs attributes are efficient.

If I need high-speed zero-copy access, I can expose memory via `mmap()` and `remap_pfn_range()`.

Bonus tip:

Mention security: "Always validate user pointers and protect shared data with synchronization."

7 Explain `probe()` and `remove()` in a platform driver.**Concept:**

They're called when a device is matched or removed.

Technical Answer:

probe() is invoked when a platform device's compatible string matches the driver's of_match_table.

Inside probe, we parse the device tree, get resources (platform_get_resource, platform_get_irq), map registers, request IRQs, and initialize hardware.

remove() is called when the device is removed or the driver is unloaded.

It performs cleanup like freeing IRQs, disabling hardware, and unregistering interfaces.

Bonus tip:

Mention devm_ functions: "I use devm APIs for automatic cleanup, which simplifies my remove() function."

8 What happens when you do cat /dev/mychar?**Concept:**

The VFS and your driver interact through file_operations.

Technical Answer:

When user does cat /dev/mychar, VFS looks up the device file's major and minor, finds the registered driver, and calls its open() and then read() methods.

Inside read(), driver copies data to user space using copy_to_user().

If no data is available, it may block on a wait queue until data is ready.

Bonus tip:

You can add: "If O_NONBLOCK is set, I return -EAGAIN instead of sleeping."

9 How do you handle interrupts inside your driver?**Concept:**

You register an ISR and manage deferred processing.

Technical Answer:

I use `request_irq()` or `devm_request_threaded_irq()` to register the ISR.

In the top-half handler, I perform minimal work like reading status registers and then defer heavy tasks using workqueues or tasklets.

I always return `IRQ_HANDLED` to indicate it's my interrupt.

During cleanup, I call `free_irq()` in `remove()`.

Bonus tip:

Mention `IRQ_WAKE_THREAD`: "If I need to sleep in ISR, I use threaded IRQ with `IRQF_ONESHOT`."

10 Difference between Sysfs and Procfs? When to use each?**Concept:**

Both expose kernel info to user space, but for different purposes.

Technical Answer:

`sysfs` is used to expose device attributes as files under `/sys` and is part of the kernel object model. Each attribute represents a specific parameter or state.

`procfs` is used for general kernel or process information, often under `/proc`. It's not specific to devices.

For device drivers, we prefer `sysfs` because it's structured and automatically linked with the device model.

Bonus tip:

Add: "For debugging-only data, I use `debugfs` instead of `procfs` since it can be disabled easily in production."

Q: If your driver is hanging after `request_irq`, how do you debug it?

- Check if interrupt line is correct (verify DT interrupts property).
- Confirm with `/proc/interrupts` whether the IRQ is firing.

- Add printk or ftrace to check if ISR runs.
- If kernel panic or oops, capture logs via serial or net console.
- Use `cat /proc/irq/<n>/` and enable `echo 1 > /proc/irq/<n>/smp_affinity_list` to debug affinity.

Impress Tip:

Add: "I also use `request_threaded_irq()` to separate top-half and bottom-half processing when debugging complex interrupts."

BSP & Yocto Build System**1 — High-level BSP / Boot flow (what runs in which stage)**

1. **ROM/ROM-bootloader** (on SoC) — first code executed from on-chip ROM/boot ROM.
2. **Primary bootloader (SPL / TPL)** — minimal loader to initialize RAM, clocks, DRAM. (Often called SPL in U-Boot world.)
3. **U-Boot (full)** — init device interfaces, provide console, env, boot commands, load kernel/initramfs/dtb.
4. **Kernel** — decompress, init device tree, probe drivers, mount rootfs.
5. **Init (systemd / init)** — userspace starts services and runs application.

Interview tip: mention SPL for boards with DRAM requiring early init and explain difference between SPL and full U-Boot.

2 — Typical BSP bring-up checklist (step by step)

1. **Serial console ready** on board (115200 8N1 commonly) — essential for kernel logs.
2. **SPL / U-Boot** compile & flash; verify U-Boot prompt present.

3. **Device tree (DTS):** create/edit .dts for SoC board, cover memory, clocks, pinctrl, peripherals.
4. **Kernel build & DTB:** build kernel with menuconfig and generate .dtb.
5. **Rootfs:** build minimal rootfs (BusyBox, or Yocto image) and copy to storage.
6. **Boot sequence:** use U-Boot to load DTB, kernel (zImage/Image), and rootfs (initramfs or block device).
7. **Driver probing:** check dmesg for device binding and errors.
8. **Enable peripherals:** test GPIO, I2C, SPI, UART, Ethernet, etc.
9. **Iterate:** fix DT properties and kernel config, rebuild.

Concrete interview phrase: “I always bring up serial console first, then U-Boot, then kernel + rootfs, and iterate on DT and kernel config until devices probe.”

3 — Device Tree: structure & common gotchas

- **Structure:** root node / { model = "..."; compatible = "..."; chosen { ... }; }; then SoC node with child nodes for peripherals.
- **Key properties:**
 - compatible — used to match drivers (of_match_table).
 - reg — address/size for MMIO.
 - interrupts — IRQ numbers / interrupt specifier.
 - pinctrl-0, pinctrl-names — pin mux configuration.
 - clocks and clock-names — clock provider references.
 - gpios — GPIO references for pins.

- status = "ok" / "disabled".
- **Common mistakes**
 - Wrong compatible string → driver not matched.
 - Incorrect reg values or missing ranges.
 - Wrong interrupt specifier (GIC uses different tuples).
 - Missing pinctrl or clock bindings → device fails to probe.
- **Debugging**
 - dmesg | grep -iOF look for of_parse, of_match, failed to get clock, pinctrl, regmap errors.
 - Use dtc -I dtb -O dts <file.dtb> to inspect the DTB on host.
 - Add #address-cells and #size-cells correct values on bus nodes.

Example snippet (I2C device):

```
&i2c1 {
    status = "okay";
    clock-frequency = <100000>;
    rtc@68 {
        compatible = "rohm,rv8803";
        reg = <0x68>;
        interrupt-parent = <&gpio1>;
        interrupts = <10 IRQ_TYPE_LEVEL_LOW>;
```

```
};  
};
```

Interview tip: explain how driver uses `of_property_read_u32()` and `of_get_named_gpio()`.

4 — U-Boot: key features & commands to know

- **Important build steps**

- Configure board defconfig in U-Boot make `<board>_defconfig`
- make `CROSS_COMPILE=arm-linux-gnueabihf-`
- Generate FIT images (`mkimage -f fit.its`) when using combined kernel+dtb+initrd.

- **Useful runtime commands**

- `printenv / setenv / saveenv`
- `mmc rescan, mmc list, ext4ls, fatload mmc 0:1 0x82000000 zImage`
- `bootz / bootm / booti (zImage/legacy/uImage/ARM64)`
- `fdt addr / fdt resize / fdt apply`
- `tftpboot`

- **Tips**

- Use `setenv bootargs` to change kernel command line (e.g., `console, root=`).
- Use `earlycon` or `earlyprintk` (kernel side) to get early kernel prints.

Interview tip: mention env variables `bootcmd`, `bootargs`, and how you use `boot.scr` or `extlinux.cfg`.

5 — Kernel configuration & building for board

- **Kernel config basics**

- `make ARCH=arm CROSS_COMPILE=... menuconfig`
- Key choices: enable Device Tree Support, required drivers (I2C, SPI, MTD, eMMC), networking, filesystems.
- Enable `CONFIG_DEBUG_INFO` only if you need symbols for `kgdb` (but increases size).

- **DT integration**

- Place DTS in `arch/arm/boot/dts/` or supply dtbs via kernel Makefile `dtbs` target.

- **Modules vs built-in**

- Put core drivers as built-in if used early (e.g., storage driver needed to mount rootfs).
- Modules useful for later loading — but if rootfs mounts over a device driven by module, kernel will fail to mount.

- **Make targets**

- `make zImage dtbs modules`
- `make modules_install` used for host staging

Interview tip: say you decide between built-in and module based on whether driver is needed before rootfs mount (e.g., eMMC driver must be built-in if rootfs on eMMC).

6 — Yocto: structure & important concepts

- **Layers:** collections of recipes and configuration. Common ones:

- poky (reference), meta-yocto, meta-openembedded, meta-<vendor>, meta-<your-bsp>.
- **Recipes** (.bb): instructions to build a package.
- **bbappend** (.bbappend): append/override a recipe from another layer.
- **Conf files**:
 - bblayers.conf — list of layers used by build.
 - local.conf — local build configs (MACHINE, parallelism, SDK settings).
- **Image recipes** (core-image-minimal, core-image-base, core-image-sato) — define package sets.
- **Classes** (.bbclass) — common recipe behavior, e.g., kernel.bbclass.
- **Devicetree in Yocto**
 - Kernel recipes handle DTBs — often you add custom .dts to kernel recipe SRC_URI or provide a devicetree.bbappend.
- **Custom BSP layer**
 - Create meta-myboard with conf/machine/<myboard>.conf and your recipes under recipes-*.
 - Set MACHINE in local.conf to your machine.
- **BitBake**
 - Key commands: bitbake <image> to build images, bitbake -c cleansstate <recipe>, bitbake -c devshell <recipe>, bitbake -s for overview.

Example machine conf:

```
# conf/machine/myboard.conf
```

```
require conf/machine/include/tune-arm.inc
```

```
MACHINE_FEATURES += "wifi bluetooth"
```

```
PREFERRED_PROVIDER_virtual/kernel = "linux-myboard"
```

```
MACHINE_ARCH = "arm"
```

Interview tip: mention meta-custom layer for company BSP and meta-yocto for common recipes.

7 — Adding a custom kernel / DTS to Yocto (practical recipe steps)

1. **Create kernel recipe** recipes-kernel/linux/linux-myboard_<version>.bb or use meta-myboard/recipes-kernel/linux/ and provide SRC_URI to your kernel git.
2. **Provide DTBs:**
 - Add SRC_URI += "file://myboard.dts file://myboard.dtsi"
 - Use KERNEL_DEVICETREE or DEVICE_TREE variable in kernel recipe to list DTBs.
3. **Kernel config**
 - Add defconfig or fragment files; use KERNEL_DEFCONFIG or FILESEXTRAPATHS.
4. **bbappend to modify existing kernel:**

```
# linux-yocto_%.bbappend
```

```
FILESEXTRAPATHS_prepend := "${THISDIR}/files:"
```

```
SRC_URI += "file://myboard.dts file://my-defconfig"
```

```
KERNEL_DEVICETREE += "myboard.dtb"
```

5. **Create machine config** (meta-myboard/conf/machine/myboard.conf) and ensure layer priority in bblayers.conf.

Concrete interview phrase: “I usually supply kernel patches and devicetree as files in the layer and use a .bbappend to append them to the main kernel recipe so upstream changes are easier to track.”

8 — Building image and deploying to target

- **Common build:**
 - source oe-init-build-env
 - Edit conf/local.conf: set MACHINE, DL_DIR, SSTATE_DIR, BB_NUMBER_THREADS.
 - bitbake core-image-minimal or core-image-base
- **Deploy:**
 - For SD card: dd if=tmp/deploy/images/<machine>/core-image-minimal-<machine>.wic of=/dev/sdX bs=4M && sync
 - For UBI/MTD: use ubinize and flash_erase/nandwrite tools or mtd tools as per device.
 - For flashing eMMC: use U-Boot fatload + dd or vendor flashing tool.
- **SDK**
 - bitbake -c populate_sdk core-image-minimal to generate cross-toolchain SDK.

Interview tip: explain wic image layouts and how you create partitions using wks file.

9 — Debugging boot issues — checklist & commands

1. **No serial logs at all**
 - Ensure serial pins/baud are correct; check UART enabled in U-Boot; check console= kernel argument.

2. U-Boot doesn't start

- Check if SPL flashed properly; check BootROM logs if available; reflash U-Boot.

3. Kernel panics / hangs early

- Enable earlyprintk / loglevel=8 and earlycon in bootargs.
- Use CONFIG_DEBUG_LL or printk in kernel.
- Check DT: missing memory node or wrong reg.

4. Driver not binding

- dmesg | grep -i <driver-prefix>; check of_match messages; ensure compatible matches.

5. Rootfs not found

- Check bootargs root= parameter; is device driver for storage built-in? check dmesg for eMMC/SD errors.

6. Netconsole / kgdb

- Use netconsole to forward kernel logs over Ethernet if serial not available.
- Use kgdb (requires kernel debug support and serial/ethernet).

7. Use perf / ftrace for performance / hangs.**8. Verify DTB used** via dmesg - kernel prints which DTB was used.

Useful commands to run on host and target:

- Host: dtc -I dtb -O dts <file.dtb>
- Target: cat /proc/interrupts, cat /proc/devices, lsmod, dmesg | tail -n 200
- U-Boot: env print, mmc info, fatls mmc 0:1

Interview tip: show structured approach: check early serial → bootloader → kernel → dt → rootfs.

10 — Yocto advanced practices for BSP maintenance

- **Layer structure:** keep meta-myboard for machine + kernel recipes, meta-myboard-boot for bootloader/U-Boot if needed.
- **Use bbappend not fork:** prefer .bbappend for small changes and maintainability.
- **SSTATE & DL cache:** configure shared SSTATE_DIR for faster iterative builds.
- **CI:** run bitbake -k in CI and test images with QEMU (if possible) or hardware farms.
- **Version pinning:** pin kernel recipe hash in layers or your layer's README to avoid unexpected upstream changes.

Interview tip: mention you automate bitbake -c cleansstate selectively to save time.

11 — Common interview questions (BSP & Yocto) + model answers

Q1 — *How do you add a new board to Yocto?*

A: Create a machine config (meta-myboard/conf/machine/myboard.conf), add kernel recipe (or bbappend) providing DTS and defconfig, add U-Boot recipe if needed, update bblayers.conf and build image for MACHINE=myboard. Use bitbake core-image-minimal.

Q2 — *When do you build kernel as module vs built-in?*

A: If the driver is required to mount rootfs (storage, driver for root device), build it **built-in**. If driver can be loaded later and not needed early, use module to reduce footprint and allow updates.

Q3 — *How do you debug device tree issues?*

A: Use dtc to decompile DTB, check dmesg for of_parse errors, ensure compatible strings match of_match_table in driver, verify clocks/pinctrl/gpio properties, and enable debug prints in driver.

Q4 — *How to bring up network on new board?*

A: Ensure network PHY and MAC clocks and pinctrl configured in DT, ensure PHY phydev is present, enable driver in kernel config (built-in if early network needed), check ifconfig -a/ip link, and use ethtool or mii-tool.

Q5 — *What is a wks file and why use it?*

A: .wks is a WIC kickstart file describing partitions and content for wic image generation. It's used to create partitioned images like SD card images with boot and rootfs partitions.

Q6 — *How to add device tree overlays?*

A: For boards supporting overlays (e.g., Raspberry Pi style), you add overlays compiled from .dts to the boot partition and update bootloader config to apply them. On U-Boot, you can apply overlays using fdt apply where supported. But for most embedded boards, create a full DTS per board variant.

Q7 — *How do you test a kernel change without full Yocto rebuild?*

A: Build only the kernel recipe (bitbake -c compile virtual/kernel or bitbake -c compile linux-myboard), deploy zImage and DTB to boot medium, and reboot. Use sstate-cache to speed up repeated builds.

12 — Practical snippets & commands (copyable)

Build kernel and DTB locally:

```
export ARCH=arm
```

```
export CROSS_COMPILE=arm-linux-gnueabi-
```

```
make myboard_defconfig
```

```
make zImage dtbs -j8
```

```
# dtb at arch/arm/boot/dts/myboard.dtb
```

Add DT to Yocto kernel recipe via bbappend:

```

# linux-yocto_%.bbappend

FILESEXTRAPATHS_prepend := "${THISDIR}/files:"

SRC_URI += "file://myboard.dts \
            file://myboard.dtsi \
            file://my_defconfig"

KERNEL_DEVICETREE += "myboard.dtb"

Flash wic image to SD (host):

sudo dd if=tmp/deploy/images/myboard/core-image-minimal-myboard.wic of=/dev/sdX bs=4M conv=fsync

U-Boot flow (load and boot):

# load kernel (zImage) at 0x82000000, dtb at 0x83000000, initramfs at 0x84000000

fatload mmc 0:1 0x82000000 zImage

fatload mmc 0:1 0x83000000 myboard.dtb

bootz 0x82000000 - 0x83000000

Enable netconsole (kernel cmdline):

netconsole=@<src-ip>:<src-port>@<dev>:<dst-ip>:<dst-port>

# e.g. netconsole=@192.168.1.10/6666@eth0/192.168.1.2/6666

```

13 — Practical interview/real brings: a short plan you can say in interview

1. Get serial console working — confirm U-Boot prompt.

2. Confirm memory (DT memory node) — kernel must know RAM.
3. Boot a known-good kernel/rootfs image (from reference board) on new board — if works, hardware is mostly correct.
4. If fails, check U-Boot env and bootargs; enable earlyprintk.
5. Iterate on DT: start with a minimal DT and incrementally add peripherals.
6. Ensure clock/pinctrl and power domains correct — these are frequent causes.
7. Once kernel boots, test storage, network and other peripherals and make drivers built-in as needed.
8. Integrate artifacts into Yocto (machine conf, kernel bbappend, rootfs packages).
9. Create reproducible build and SDK for development.

Yocto Build System.

1. Yocto Basics

Yocto is not a Linux distro — it's a **build framework** used to create *custom embedded Linux distributions*.

◆ Components:

Component	Description
Poky	The reference distribution of Yocto (includes BitBake + core metadata). Think of it as a base project.
BitBake	The build engine — it parses recipes and executes tasks to compile and package software.

Component	Description
Layers	A collection of metadata (recipes, configuration, classes). Each layer adds specific functionality (e.g., BSP, UI, tools).
Recipes (.bb)	Instructions on <i>how to fetch, configure, compile, and install</i> a package.
Metadata	Information that drives the build — includes .bb, .bbclass, .bbappend, and configuration files.

2. Metadata Hierarchy

Yocto uses a well-structured metadata system.

File Type	Purpose
conf/	Contains layer configuration files (like layer.conf, local.conf, bblayers.conf).
.bb	Recipe file: defines how to build one package (e.g., busybox_1.35.bb).
.bbappend	Used to extend or modify an existing recipe (e.g., add patches, change variables).
.bbclass	Shared class files (common logic reused by many recipes, like autotools.bbclass).

Example:

meta-myproject/

└─ conf/layer.conf

```

└── recipes-core/
    ├── busybox/
    │   ├── busybox_1.35.bbappend
    │   └── myapp/
    │       └── myapp_1.0.bb

```

3. Custom Image Creation

An **image** recipe defines what packages go into the final root filesystem.

Example:

```
recipes-core/images/core-image-mycustom.bb
```

```
SUMMARY = "My Custom Embedded Linux Image"
```

```
LICENSE = "MIT"
```

```
inherit core-image
```

```
IMAGE_INSTALL += "busybox openssh nano"
```

```
IMAGE_FEATURES += "ssh-server-dropbear"
```

Then build it:

```
bitbake core-image-mycustom
```

4. Adding Layers and Recipes

Add new layer:

```
bitbake-layers create-layer ../meta-myproject
```

```
bitbake-layers add-layer ../meta-myproject
```

Add new recipe:

```
cd meta-myproject/recipes-apps/
```

```
mkdir myapp && cd myapp
```

```
bitbake-layers create-recipe myapp
```

5. Modifying Existing Recipes (.bbappend)

If you want to *modify or patch* an existing recipe (without touching the original):

Example:

File: meta-myproject/recipes-core/busybox/busybox_1.35.bbappend

```
SRC_URI += "file://my_patch.patch"
```

```
do_install_append() {
```

```
    install -m 0755 ${WORKDIR}/my_script.sh ${D}${bindir}/
```

```
}
```

6. local.conf vs bblayers.conf

File	Purpose
------	---------

local.conf	User-level build settings (machine, image type, parallel threads, etc). Located in build/conf/.
-------------------	---

bblayers.conf	Defines which layers are included in the build. (path to each meta-layer).
----------------------	--

Example:

```
# local.conf
```

```
MACHINE ?= "beaglebone"
```

```
DISTRO ?= "poky"
```

```
IMAGE_INSTALL:append = " vim"
```

```
# bblayers.conf
```

```
BBLAYERS += " \
```

```
  /home/user/poky/meta \
```

```
  /home/user/meta-openembedded/meta-oe \
```

```
  /home/user/meta-myproject \
```

```
"
```

7. BitBake Build Flow

1. **bitbake core-image-minimal**
 2. BitBake parses all metadata (conf, .bb, .bbclass, .bbappend)
 3. It builds the **dependency graph**
 4. Executes *tasks* for each recipe:
fetch → unpack → patch → configure → compile → install → package → rootfs
 5. Generates output in:
 6. tmp/ deploy/ images/ <machine> /
 7. tmp/ deploy/ ipk/
-

8. Deploying Image to Target

Once build completes:

```
cd tmp/ deploy/ images/ beaglebone/
```

```
ls
```

```
# core-image-minimal-beaglebone.ext4
```

```
# MLO, u-boot.img, zImage, dtb
```

```
# Flash using dd or bmaptool
```

```
sudo dd if=core-image-minimal-beaglebone.ext4 of=/dev/sdX bs=4M
```

9. Debugging Yocto Build Failures

Common methods:

- Use verbose logs:
- `bitbake <recipe> -c compile -f -v`
- Inspect log files:
- `tmp/work/<machine>/<recipe>/temp/log.do_compile`
- Clean and rebuild:
- `bitbake -c cleanall <recipe>`
- `bitbake <recipe>`
- Dependency tracking:
- `bitbake -g <image>`
- `cat pn-depends.dot | dot -Tpng > deps.png`

1 What's the difference between recipe and layer?

Aspect	Recipe	Layer
Definition	Single build instruction for one package	Collection of related recipes, configs, and classes
File	.bb, .bbappend	Directory (meta-*) with conf/ and recipes
Purpose	Defines <i>how</i> to build software	Organizes metadata logically (BSP, UI, app, etc.)

Aspect	Recipe	Layer
Example	busybox_1.35.bb	meta-openembedded, meta-myproject

2 How do you add a custom kernel in Yocto?

Steps:

1. Create or clone kernel source layer (e.g., meta-my-kernel)
 2. Add kernel recipe:
 3. recipes-kernel/linux/linux-myproject_6.1.bb
 4. require recipes-kernel/linux/linux-yocto.inc
 5. SRC_URI = "git://github.com/myorg/linux.git;branch=mybranch"
 6. Update local.conf:
 7. PREFERRED_PROVIDER_virtual/kernel = "linux-myproject"
 8. Add layer & build image.
-

3 How to create your own image with specific packages?

1. Create new image recipe:
2. meta-myproject/recipes-core/images/core-image-mycustom.bb
3. Add required packages:

4. `IMAGE_INSTALL += "bash openssh vim myapp"`
5. Build:
6. `bitbake core-image-mycustom`

Topic	Key Answer
What is Poky?	Poky is the reference distribution of Yocto (includes BitBake and core layers).
What are Layers?	Modular metadata collections (e.g., meta-ti for TI boards).
How to add custom kernel?	Create your own kernel recipe & set <code>PREFERRED_PROVIDER_virtual/kernel</code> .
What is .bbappend?	Used to modify or extend existing recipes.
local.conf vs bblayers.conf	<code>local.conf</code> = user build settings; <code>bblayers.conf</code> = active layers.
How to add new app?	Create recipe in your meta-layer's <code>recipes-apps/</code> .
How to debug?	Use verbose mode, inspect <code>log.do_compile</code> , and clean builds.

Kernel Internals

1 Kernel Architecture Overview

The Linux kernel has four major subsystems:

Subsystem	Description
Process Management	Handles tasks (processes, threads), scheduling, and context switching.
Memory Management	Manages virtual memory, physical pages, kmalloc/vmalloc, and paging.
File System / I/O	Implements VFS (Virtual File System) and drivers for block/char devices.
Device Drivers & Networking	Handle hardware interfaces, network stacks, and protocols.

Kernel Space vs User Space

- **User Space** → Applications run here, isolated from hardware.
 - **Kernel Space** → Privileged mode; direct hardware access.
 - System calls are the bridge between both.
-

2 Process Management

◆ Processes & Threads

Each **process** has a `task_struct` (in `include/linux/sched.h`) that holds info like PID, scheduling policy, CPU time, etc. Both **user threads** and **kernel threads** use the same `task_struct`.

◆ Scheduler

- Uses **Completely Fair Scheduler (CFS)** for normal tasks.
- **Real-time schedulers:** SCHED_FIFO, SCHED_RR.
- Every task has a **priority** and **time slice**.

◆ Context Switching

When CPU switches from one process to another:

1. Save current task's registers and state (in task_struct).
2. Load next task's saved state.
3. Update mm_struct (memory map) and CPU context.

🔧 Done by:

switch_to(prev, next) → Assembly-level context change.

3 Memory Management

kmalloc()

- Allocates **physically contiguous** memory.
- Used for DMA buffers or small kernel allocations.
- Wrapper over the **slab allocator**.

vmalloc()

- Allocates **virtually contiguous** memory (not physically contiguous).

- Suitable for large memory blocks not used for DMA.

Slab Allocator

- Caches frequently used kernel objects (e.g., task_struct).
 - Reduces allocation overhead using object “slabs” from a common pool.
-

System Call Mechanism

System calls are the interface between user space and kernel space.

Flow Example (x86_64):

1. User app calls read().
2. libc converts to syscall instruction.
3. CPU jumps to **syscall handler** in kernel (entry_SYSCALL_64).
4. Handler uses the **syscall table** (sys_call_table[]) to find the function pointer.
5. Executes kernel function like sys_read().
6. Returns result via CPU register.

Syscall number is stored in rax register (x86) or x8 (ARM64).

Interrupt Handling Flow

Hardware to Kernel Flow:

1. **Hardware IRQ** → Interrupt Controller (GIC on ARM).

2. CPU switches to **IRQ mode** and jumps to the **vector table**.
3. Calls `do_IRQ()`.
4. Calls the registered **ISR (Interrupt Service Routine)** for that IRQ.
5. ISR handles minimal work → schedules a **bottom half** using:
 - **Tasklets, Workqueues, or SoftIRQs.**
6. Return from interrupt → normal execution resumes.

 Example (GPIO IRQ):

```
request_irq(irq_number, my_isr, IRQF_TRIGGER_RISING, "mydev", NULL);
```

Workqueues, Timers, and Kthreads

Mechanism	Purpose
Workqueue	Defers work to be executed in process context (can sleep).
Timer	Executes function after a timeout (<code>mod_timer()</code>).
Kernel Thread (kthread)	Persistent thread in kernel for background work.

Example:

```
DECLARE_WORK(my_work, my_func);
schedule_work(&my_work); // executes in process context
```

```
struct task_struct *task = kthread_run(my_thread_func, NULL, "my_kthread");
```

7 Synchronization Primitives

Primitive	Used In	Can Sleep?	Description
Spinlock	IRQ or atomic context	✗ No	Busy-waits; fast for short sections.
Mutex	Process context	✓ Yes	Sleeps if locked; for longer critical sections.
Semaphore	Process context	✓ Yes	Can be used for resource counting.
Atomic_t	Anywhere	✗ No	Atomic integer operations.

Always use **spinlocks** inside ISRs or atomic context (cannot sleep).

Use **mutex/semaphore** in normal process context.

8 Kernel Modules and Linking

Kernel modules are dynamically loadable objects (.ko files).

- Build via make modules.
- Load: insmod mymodule.ko
- Remove: rmmod mymodule
- Info: modinfo mymodule.ko

Example Makefile

```
obj-m += mydriver.o
```

```
all:
```

```
make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
```

9 Exported Symbols and Module Dependencies

- Symbols exported with:
- `EXPORT_SYMBOL(my_function);`

or

```
EXPORT_SYMBOL_GPL(my_function);
```

- Other modules can use it by declaring `extern`.

View dependencies:

```
modinfo mydriver.ko
```

1 Explain the difference between Kernel Thread and User Thread

Feature	Kernel Thread	User Thread
Execution Mode	Runs in kernel space only	Runs in user space , uses syscalls to access kernel
Creation	<code>kthread_create()</code> or <code>kthread_run()</code>	Created via <code>pthread</code> s or <code>fork()</code>

Feature	Kernel Thread	User Thread
Scheduling	Managed by kernel scheduler directly	Scheduled like normal process
Access to Hardware	Direct access (can use kernel APIs)	No direct hardware access
Example	ksoftirqd, kswapd	User applications (Firefox, bash)

Kernel threads are system daemons doing kernel-level work; user threads are for applications.

2 How does Context Switching happen in Linux?

- Triggered by **scheduler** or **interrupt**.
- Steps:
 1. Save current task's CPU context (registers, stack pointer).
 2. Load next task's context from its task_struct.
 3. Switch memory mapping (if user task).
 4. Update CPU state → jump to next process.

Function: `switch_to(prev, next)` (arch-specific ASM).

Note: Context switching is expensive → kernel minimizes it using efficient scheduling.

3 Explain Kernel Preemption and Its Impact

Type	Description
Non-Preemptible Kernel	Once kernel starts executing, cannot be preempted until return to user mode or explicit schedule point.
Preemptible Kernel	Kernel code can be preempted by higher-priority tasks — improves real-time response.
CONFIG_PREEMPT_RT	Real-time patchset for deterministic latency.

Impact:

- Increases responsiveness (good for RT systems).
- Slightly increases overhead (more context switches).
- Controlled by CONFIG_PREEMPT in kernel config.

Quick Revision Summary

Concept	Keyword / Command
Scheduler	CFS, switch_to()
Memory Allocators	kmalloc (physically cont.), vmalloc (virtually cont.)
System Call	sys_call_table, entry_SYSCALL_64
Interrupt Flow	IRQ → do_IRQ → ISR → bottom half
Deferred Work	Workqueues, tasklets, timers
Sync Primitives	spinlock, mutex, semaphore, atomic

Concept	Keyword / Command
Kernel Threads	kthread_run()
Module Linking	EXPORT_SYMBOL(), insmod, modinfo
Preemption	CONFIG_PREEMPT, improves latency

E. Kernel Debugging & Analysis — Deep Core Explanation

1. Kernel Logs (dmesg, printk levels)

What happens

- The Linux kernel uses an internal circular buffer called the kernel log buffer (/proc/kmsg or /dev/kmsg) to store messages from the kernel and drivers.
- printk() is the kernel-space equivalent of printf() in user space.
- Messages are visible via dmesg command or journalctl -k (in systemd-based systems).

printk syntax

```
printk(KERN_INFO "Device initialized successfully\n");
```

Log levels (importance priority)

Level	Macro Name	Description	Typical Use
0	KERN_EMERG	Emergency (system unusable)	Hardware failure
1	KERN_ALERT	Immediate action required	Disk failure
2	KERN_CRIT	Critical conditions	Kernel panic
3	KERN_ERR	Error conditions	Driver failure
4	KERN_WARNING	Warning	Resource limit approaching
5	KERN_NOTICE	Normal but significant	Important info
6	KERN_INFO	Informational	Device probe success
7	KERN_DEBUG	Debug-level messages	Debugging only

Commands

`dmesg | less`

`dmesg -T` # show timestamps in human-readable form

`dmesg --level=err,warn` # filter error and warning messages

2. Oops vs Panic

Kernel Oops

- An Oops occurs when the kernel detects an illegal operation (e.g., NULL pointer dereference, invalid memory access) but recovers gracefully.
- It kills the offending process, but the system may still be *partially functional*.
- Logged in /var/log/kern.log or dmesg.

Example cause: accessing NULL pointer in driver probe()

Example message:

Unable to handle kernel NULL pointer dereference at virtual address 00000000

CPU: 0 PID: 1234 Comm: my_driver

PC is at my_driver_probe+0x10/0x30 [my_driver]

Kernel Panic

- More severe. System state is corrupted, so kernel halts execution to prevent further damage.
- Panic is triggered by panic() call inside the kernel or by severe errors (e.g., unrecoverable Oops, filesystem corruption, memory corruption).

Example log:

Kernel panic - not syncing: Fatal exception

CPU: 1 PID: 0 Comm: swapper/1 Not tainted 6.1.1

Call Trace:

dump_stack+0x9c/0xc8

panic+0x11e/0x30a

Difference summary

Feature	Oops	Panic
Severity	Recoverable	Non-recoverable
System state	Partial recovery	System halts/reboots
Trigger	Invalid memory, bad pointer	Critical unrecoverable fault
Result	Process killed	Kernel halted
Debug goal	Fix faulty driver/module	Identify crash root cause

3. Crash Dump Analysis

When a kernel panics, the system can be configured to capture memory contents (core dump) for post-mortem analysis.

Tools:

- kdump: kernel crash dump mechanism.
- makedumpfile: extracts essential info from /proc/vmcore.
- crash: interactive utility for analyzing dump files.

Flow:

1. Primary kernel runs normally.
2. On panic, kexec loads a secondary crash kernel.
3. The crash kernel saves /proc/vmcore.
4. Use crash tool to analyze:

5. `crash /usr/lib/debug/vmlinux /var/crash/vmcore`
 6. Within crash, analyze:
 - `bt` — backtrace
 - `ps` — process list
 - `kmem` — memory usage
 - `log` — kernel log buffer
-

4. kgdb / kdb / gdb

1. kgdb (Kernel GDB)

- Allows remote debugging of a running kernel over serial or Ethernet.
- Similar to user-space gdb debugging.
- Useful for breakpoints, stepping, inspecting memory/variables.

Setup:

- Enable in kernel config:
- `CONFIG_KGDB=y`
- `CONFIG_KGDB_SERIAL_CONSOLE=y`
- `CONFIG_FRAME_POINTER=y`
- Boot arguments:
- `kgdboc=ttyS0,115200`

- kgdbwait
- Connect from host:
- gdb vmlinux
- (gdb) target remote /dev/ttyS0
- (gdb) continue

2. kdb

- Built-in simple debugger (like interactive shell in kernel).
- Access via console on panic.
- Useful for checking stack, memory, processes:
- kdb> bt
- kdb> ps
- kdb> lsmod

3. gdb with qemu or crash dump

- If kernel runs in QEMU:
- qemu-system-arm -kernel zImage -S -s
- gdb vmlinux
- (gdb) target remote :1234

5. Tracing and Performance Tools

strace

- Traces system calls from a user-space process.
- Helps differentiate user-space issues vs kernel-space.

strace -p <pid>

strace ./app

perf

- Performance profiler (CPU cycles, cache misses, etc.)
- Records performance events:
- perf record -a
- perf report

ftrace

- Internal kernel tracer.
- Located at /sys/kernel/debug/tracing/.
- Can trace function calls, IRQs, schedulers, preemptions.
- echo function > /sys/kernel/debug/tracing/current_tracer
- cat /sys/kernel/debug/tracing/trace

Common tracers:

function, function_graph, irqsoff, wakeup_rt.

◆ 6. Static Analysis Tools

smatch

- Detects common kernel coding issues (e.g., uninitialized vars, null deref).
- Used during kernel build:
- `make C=1 CHECK="smatch"`

sparse

- Semantic static analyzer — finds type mismatches, missing `__user` annotations.
 - `make C=2 CHECK="sparse"`
-

◆ 7. Kernel Panic Troubleshooting Steps

Step-by-step approach

1. Collect evidence

- Serial logs (dmesg, console logs).
- Crash dump (kdump, /proc/vmcore).
- Kernel version, configuration, custom patches.

2. Identify panic signature

- Look for "Kernel panic - not syncing" or Oops.
- Analyze the *Call Trace* and *Faulting Address*.

3. Analyze call trace

- The last few entries in the backtrace show the *function chain*.
- Look for your driver/module name.
- Example:
- [`c0301210`] `my_driver_probe+0x10/0x50 [my_driver]`
- [`c0300a60`] `platform_drv_probe+0x20/0x40`
- The topmost function is the **fault origin**.

4. Reproduce with debug info

- Enable dynamic debugging: `echo 'module my_driver +p' > /sys/kernel/debug/dynamic_debug/control`
- Add more `printk(KERN_DEBUG ...)` messages.

5. Use **kgdb** or **ftrace** to trace real-time flow.

6. **Fix and re-test** with instrumented build.

Expected Interview Questions — Detailed Answers

Q1: How to identify root cause of a kernel panic?

Answer:

1. Capture logs via serial or `dmesg -T`.
2. Identify “Call Trace” section and top function causing crash.
3. Match address with symbol via:

4. `addr2line -e vmlinux 0xc0301210`
 5. Check for NULL dereferences, memory corruption, race conditions.
 6. If reproducible, use `kgdb` or `ftrace` to trace the path.
 7. Optionally use `crash /proc/vmcore vmlinux` for postmortem analysis.
-

Q2: Difference between Oops and Panic?

✓ **Oops:** recoverable kernel fault; system continues.

✓ **Panic:** unrecoverable; kernel halts.

(Explain with example — Oops from NULL deref; Panic from corrupted scheduler structures).

Q3: Methods to debug hung processes or deadlocks?

Approach:

1. Check blocked tasks: `echo w > /proc/sysrq-trigger`
 2. Use `sysrq` → `SysRq + t` to dump task states.
 3. Trace using `ftrace` or `perf sched record`.
 4. Use `lockdep (CONFIG_PROVE_LOCKING)` to detect deadlocks.
 5. Examine `ps` output in `crash` tool to see which thread is waiting for which lock.
-

Q4: How to analyze call trace in kernel logs?

Answer:

1. Each line in trace shows return address and function name.
 2. Read from bottom (earlier functions) to top (fault function).
 3. Match address using:
 4. `cat /proc/kallsyms | grep <address>`
 5. Use `addr2line` with `vmlinux` to find exact C source line.
-

Sample Kernel Panic Log – BeagleBone Black (ARM Cortex-A8)

```
[ 2.543210] my_i2c_driver: loading out-of-tree module taints kernel.
[ 2.547321] my_i2c_driver: Probing I2C device at address 0x48
[ 2.552499] Unable to handle kernel NULL pointer dereference at virtual address 00000000
[ 2.561249] pgd = c0004000
[ 2.563987] [00000000] *pgd=00000000
[ 2.567891] Internal error: Oops: 5 [#1] SMP ARM
[ 2.572678] Modules linked in: my_i2c_driver bluetooth eeprom
[ 2.578655] CPU: 0 PID: 125 Comm: kworker/0:1 Tainted: G      O   5.10.168-bbb #1
[ 2.587611] Hardware name: Generic AM335x (Flattened Device Tree)
[ 2.593821] Workqueue: events my_i2c_work_handler
[ 2.598877] PC is at my_i2c_probe+0x14/0x40 [my_i2c_driver]
```

```

[ 2.604438] LR is at i2c_device_probe+0x58/0x80
[ 2.609354] pc : [<bf000014>]   lr : [<c038a058>]   psr: 60000013
[ 2.615570] sp : c36a9f80 ip : 00000000 fp : c36a9fa4
[ 2.620825] r10: c16b8200 r9 : 00000000 r8 : bf001000
[ 2.626137] r7 : c16b8200 r6 : c16b8200 r5 : 00000000 r4 : c16b8300
[ 2.632891] r3 : 00000000 r2 : 00000000 r1 : 00000000 r0 : c16b8200
[ 2.639645] Flags: nZCv IRQs on FIQs on Mode SVC_32 ISA ARM Segment user
[ 2.647180] Control: 10c5387d Table: 83604019 DAC: 00000051
[ 2.653017] Process kworker/0:1 (pid: 125, stack limit = 0x6b8b2f49)
[ 2.659797] Stack: (0xc36a9f80 to 0xc36aa000)
[ 2.664153] 9f80: c16b8200 bf001000 c038a058 c36a9fa4 c038a058 bf000014 c36a9fa4 c038a058
[ 2.672384] 9fa0: c00093b8 c00091c0 00000000 00000000 00000000 00000000 00000000 00000000
[ 2.680615] Backtrace:
[ 2.682878] [<bf000014>] (my_i2c_probe [my_i2c_driver]) from [<c038a058>] (i2c_device_probe+0x58/0x80)
[ 2.692478] [<c038a058>] (i2c_device_probe+0x58/0x80) from [<c0378f48>] (driver_probe_device+0x10c/0x2a4)
[ 2.702437] [<c0378f48>] (driver_probe_device+0x10c/0x2a4) from [<c03774d4>] (bus_for_each_drv+0x50/0x84)
[ 2.712325] [<c03774d4>] (bus_for_each_drv+0x50/0x84) from [<c0378a60>] (device_attach+0x9c/0xd4)
[ 2.721887] [<c0378a60>] (device_attach+0x9c/0xd4) from [<c0378adc>] (__driver_attach+0x44/0x48)
[ 2.731616] [<c0378adc>] (__driver_attach+0x44/0x48) from [<c0376b88>] (bus_add_driver+0x158/0x1f4)

```

```
[ 2.741419] Code: e3500000 0a000002 e1a03000 e5902000 (e5923000)
[ 2.747674] ---[ end trace 7c8f5b93a6d3b1e2 ]---
[ 2.751898] Kernel panic - not syncing: Fatal exception in interrupt
[ 2.757954] CPU: 0 PID: 125 Comm: kworker/0:1 Not tainted 5.10.168-bbb #1
[ 2.764923] Backtrace:
[ 2.767428] [<c000e140>] (dump_backtrace) from [<c000e420>] (show_stack+0x20/0x24)
[ 2.774914] [<c000e420>] (show_stack) from [<c000f4f4>] (panic+0x120/0x2d8)
[ 2.781826] ---[ end Kernel panic ]---
```

Step-by-Step Crash Analysis

Step 1 — Identify Crash Type

Unable to handle kernel NULL pointer dereference

Internal error: Oops: 5 [#1]

- Indicates a **NULL pointer dereference** in kernel code.
 - Initially marked as **Oops**, but since it occurred in interrupt context (Workqueue), kernel escalates to **panic** (not recoverable).
-

Step 2 — Look at the Fault Context

Workqueue: events my_i2c_work_handler

PC is at my_i2c_probe+0x14/0x40 [my_i2c_driver]

- Crash happened **inside** your driver's `my_i2c_probe()` function.
- It was running in the **workqueue context** — meaning deferred work after device detection.

Program Counter (PC) → instruction pointer (current instruction).

PC = `bf000014` indicates the faulting instruction in the driver.

Step 3 — Check Register Values

`r5 : 00000000 r4 : c16b8300`

- Register `r5` holds `NULL (0x0)`, and kernel tried to dereference it:
- `e5923000 → LDR r3, [r2]`

→ means load from address in `r2` which was `NULL`.

So:

The driver likely did something like:

```
struct my_data *data = client->dev.platform_data;
```

```
int val = data->config; // crash if data == NULL
```

Step 4 — Check Backtrace

`[<bf000014>] my_i2c_probe+0x14/0x40 [my_i2c_driver]`

`[<c038a058>] i2c_device_probe+0x58/0x80`

This shows:

1. `my_i2c_probe()` called.
2. Inside it, we accessed a NULL pointer.
3. Call came from `i2c_device_probe()`, which is part of the core I²C subsystem.

✓ **Root Cause:**

`my_i2c_probe()` didn't check if `platform_data` or device memory was valid before dereferencing.

Step 5 — Verify with `addr2line`

To find the exact C source line:

```
addr2line -e vmlinux bf000014
```

Output:

```
/home/yocto/meta-bbb/drivers/i2c/my_i2c_driver.c:128
```

Step 6 — Check for Panic Escalation

Kernel panic - not syncing: Fatal exception in interrupt

- The kernel couldn't recover because this happened inside a workqueue thread (kernel context).
 - Once the kernel detects inconsistent state, it **calls `panic()`** to avoid data corruption.
-

Step 7 — Interpret Key Lines

Line	Meaning
taints kernel	Module is not part of mainline; kernel tainted → some warnings suppressed.
pgd = c0004000	Page directory base; memory map info.
r5 : 00000000	Shows NULL pointer.
my_i2c_probe+0x14/0x40	Crash occurred 0x14 bytes from start of function.
end trace 7c8f5b93a6d3b1e2	Kernel printed the unique trace ID (useful for comparing logs).

Step 8 — How to Prevent

- Validate pointers before using:
 - `if (!client->dev.platform_data) {`
 - `dev_err(&client->dev, "No platform data!\n");`
 - `return -EINVAL;`
 - `}`
 - Use `BUG_ON()` only in debugging, not in production.
 - Enable `CONFIG_DEBUG_SLAB` and `CONFIG_DEBUG_KMEMLEAK` to detect memory misuse.
-

Step 9 — If You Had kdump Enabled

You could analyze further using:

```
crash /usr/lib/debug/vmlinux /var/crash/vmcore
```

```
crash> bt
```

```
crash> log
```

```
crash> ps
```

To see which process crashed and what memory it accessed.

Step 10 — Summary Table

Step	Task	Tool/Command
1	Capture crash logs	dmesg, serial console
2	Identify fault type	Look for Oops / panic
3	Find function	Call Trace, PC
4	Map address to source	addr2line -e vmlinux <addr>
5	Analyze registers	Look for NULL, invalid pointers
6	Cross-check build	`nm vmlinux
7	Verify fix	Rebuild driver, test again
